

homework1

September 13, 2019

NAME: Jacob Duvall

1 Homework 1

1.0.1 Objectives

- Basic numpy operations to access data
- Basic plotting of subsets of data
- Simple descriptive statistics
- Do not save work within the ml_practices folder
 - create a folder in your home directory for assignments, and copy the templates there

1.0.2 General References

- [Sci-kit Learn Breast Cancer Dataset](#)
- [Numpy Reference](#)
- [Summary of matplotlib](#)
 - Plot
 - Boxplots
 - Histograms
 - Scatter plots
 - Colormap Plots

```
[2]: import numpy as np
import pandas as pd
import matplotlib.pyplot as plt
from sklearn.datasets import load_breast_cancer
```

2 LOAD BREAST CANCER DATA SET

```
[3]: """
Load the dataset into the bc_dataset variable, by calling the
load_breast_cancer() function imported from sklearn.datasets.
Then display the bc_dataset object's list of keys. bc_dataset
is a dictionary object.
"""
```

```
bc_dataset = load_breast_cancer()
```

2.0.1 Dataset Details

The `bc_dataset` variable is a dictionary with multiple fields: * `data` : m by n numpy array of the n observed feature values, for each of the m samples

* `target` : m by 1 numpy array of samples' classification as either malignant (i.e. 0) or benign (i.e. 1)

* `target_names` : 2 by 1 numpy array of the possible tumor classifications

* `DESCR` : string containing a detailed description of the dataset

* `feature_names` : n by 1 numpy array of the names of the feature variables

* `filename` : string containing the absolute path to where the file containing all the data information is located on the local system

```
[4]: """ TODO
      Print out the description of the data, by accessing the
      'DESCR' field
      """
      print(bc_dataset.DESCR)
```

```
.. _breast_cancer_dataset:
```

Breast cancer wisconsin (diagnostic) dataset

****Data Set Characteristics:****

:Number of Instances: 569

:Number of Attributes: 30 numeric, predictive attributes and the class

:Attribute Information:

- radius (mean of distances from center to points on the perimeter)
- texture (standard deviation of gray-scale values)
- perimeter
- area
- smoothness (local variation in radius lengths)
- compactness (perimeter² / area - 1.0)
- concavity (severity of concave portions of the contour)
- concave points (number of concave portions of the contour)
- symmetry
- fractal dimension ("coastline approximation" - 1)

The mean, standard error, and "worst" or largest (mean of the three largest values) of these features were computed for each image, resulting in 30 features. For instance, field 3 is Mean Radius, field 13 is Radius SE, field 23 is Worst Radius.

- class:
 - WDBC-Malignant
 - WDBC-Benign

:Summary Statistics:

| | Min | Max |
|-------------------------------------|-------|--------|
| radius (mean): | 6.981 | 28.11 |
| texture (mean): | 9.71 | 39.28 |
| perimeter (mean): | 43.79 | 188.5 |
| area (mean): | 143.5 | 2501.0 |
| smoothness (mean): | 0.053 | 0.163 |
| compactness (mean): | 0.019 | 0.345 |
| concavity (mean): | 0.0 | 0.427 |
| concave points (mean): | 0.0 | 0.201 |
| symmetry (mean): | 0.106 | 0.304 |
| fractal dimension (mean): | 0.05 | 0.097 |
| radius (standard error): | 0.112 | 2.873 |
| texture (standard error): | 0.36 | 4.885 |
| perimeter (standard error): | 0.757 | 21.98 |
| area (standard error): | 6.802 | 542.2 |
| smoothness (standard error): | 0.002 | 0.031 |
| compactness (standard error): | 0.002 | 0.135 |
| concavity (standard error): | 0.0 | 0.396 |
| concave points (standard error): | 0.0 | 0.053 |
| symmetry (standard error): | 0.008 | 0.079 |
| fractal dimension (standard error): | 0.001 | 0.03 |
| radius (worst): | 7.93 | 36.04 |
| texture (worst): | 12.02 | 49.54 |
| perimeter (worst): | 50.41 | 251.2 |
| area (worst): | 185.2 | 4254.0 |
| smoothness (worst): | 0.071 | 0.223 |
| compactness (worst): | 0.027 | 1.058 |
| concavity (worst): | 0.0 | 1.252 |
| concave points (worst): | 0.0 | 0.291 |
| symmetry (worst): | 0.156 | 0.664 |
| fractal dimension (worst): | 0.055 | 0.208 |

:Missing Attribute Values: None

:Class Distribution: 212 - Malignant, 357 - Benign

:Creator: Dr. William H. Wolberg, W. Nick Street, Olvi L. Mangasarian

:Donor: Nick Street

:Date: November, 1995

This is a copy of UCI ML Breast Cancer Wisconsin (Diagnostic) datasets.
<https://goo.gl/U2Uwz2>

Features are computed from a digitized image of a fine needle aspirate (FNA) of a breast mass. They describe characteristics of the cell nuclei present in the image.

Separating plane described above was obtained using Multisurface Method-Tree (MSM-T) [K. P. Bennett, "Decision Tree Construction Via Linear Programming." Proceedings of the 4th Midwest Artificial Intelligence and Cognitive Science Society, pp. 97-101, 1992], a classification method which uses linear programming to construct a decision tree. Relevant features were selected using an exhaustive search in the space of 1-4 features and 1-3 separating planes.

The actual linear program used to obtain the separating plane in the 3-dimensional space is that described in:
[K. P. Bennett and O. L. Mangasarian: "Robust Linear Programming Discrimination of Two Linearly Inseparable Sets", Optimization Methods and Software 1, 1992, 23-34].

This database is also available through the UW CS ftp server:

```
ftp ftp.cs.wisc.edu
cd math-prog/cpo-dataset/machine-learn/WDBC/
```

.. topic:: References

- W.N. Street, W.H. Wolberg and O.L. Mangasarian. Nuclear feature extraction for breast tumor diagnosis. IS&T/SPIE 1993 International Symposium on Electronic Imaging: Science and Technology, volume 1905, pages 861-870, San Jose, CA, 1993.
- O.L. Mangasarian, W.N. Street and W.H. Wolberg. Breast cancer diagnosis and prognosis via linear programming. Operations Research, 43(4), pages 570-577, July-August 1995.
- W.H. Wolberg, W.N. Street, and O.L. Mangasarian. Machine learning techniques to diagnose breast cancer from fine-needle aspirates. Cancer Letters 77 (1994) 163-171.

2.1 SETUP USEFUL VARIABLES

```
[75]: """
Store the names of the features and the names
of the target classes, into the variables
feature_names and target_names respectively.
"""

feature_names = bc_dataset.feature_names
target_names = bc_dataset.target_names

""" TODO
Print the list of feature names and target names
"""

print(feature_names)
print(target_names)
```

```
['mean radius' 'mean texture' 'mean perimeter' 'mean area'
 'mean smoothness' 'mean compactness' 'mean concavity'
 'mean concave points' 'mean symmetry' 'mean fractal dimension'
 'radius error' 'texture error' 'perimeter error' 'area error'
 'smoothness error' 'compactness error' 'concavity error'
 'concave points error' 'symmetry error' 'fractal dimension error'
 'worst radius' 'worst texture' 'worst perimeter' 'worst area'
 'worst smoothness' 'worst compactness' 'worst concavity'
 'worst concave points' 'worst symmetry' 'worst fractal dimension']
['malignant' 'benign']
```

```
[74]: """
Create variables for the feature and target data
The X variable is a numpy array containing the data measured
for each feature for each sample. Each column of X is a
different feature for all the samples. Each row of X is a
different sample with all its features.
The y variable is a numpy array containing the classification
for each sample. A sample tumor is either Benign or Malignant.
"""

X = bc_dataset['data']
y = bc_dataset['target']

""" TODO
Print the dimensions of the X and y variables respectively
"""

print(X.shape)
print(y.shape)
```

```
(569, 30)
(569,)
```

```
[7]: """
      Store the number of samples and the number of features, by
      accessing the values from the shape of X
      """
      nsamples = X.shape[0]
      nfeatures = X.shape[1]

      """ TODO
      Print the print the number of samples and numberof features respectively
      """
      print(nsamples)
      print(nfeatures)
```

569

30

2.2 SELECT SUBSET OF FEATURES

Not all available data is necessary or useful for making predictions and classifying observations. There are numerous feature selection algorithms that exist, which will be discussed in more detail later within the course. For now we are going to arbitrarily select mean radius, mean area, mean concavity, and mean symmetry as our predictor variables. We will not yet be performing any predictions in this assignment; rather this term is used to conveniently distinguish this subset of features from the full set of features.

```
[8]: """
      Feature Column Indices
      The values observed for each feature resides within a particular
      column of the feature matrix, X. For example, column 0 contains the
      values of the mean radius for each observation, the column at index
      3 contains the values for the mean area, and so on.
      """
      mean_radius_idx = 0
      mean_area_idx = 3
      mean_concavity_idx = 6
      mean_symmetry_idx = 8

      """
      Create a list of the select subset of features
      """
      predictors = [mean_radius_idx, mean_area_idx, mean_concavity_idx,
                    ↪mean_symmetry_idx]

      """
      Create a variable, storing the number of predictors
      """
      npredictors = len(predictors)
```

```

"""
Create a list of corresponding names for the selected set of features.
This is conveniently done using list comprehension
"""
pred_names = [bc_dataset['feature_names'][index] for index in predictors]
"""
Print the list of predictor names
"""
print("Predictor Feature Names:")
print(pred_names)

```

Predictor Feature Names:


['mean radius', 'mean area', 'mean concavity', 'mean symmetry']



2.3 BASIC HISTOGRAMS OF FEATURES


```

[69]: """ TODO
HISTOGRAMS OF THE CHOSEN PREDICTOR FEATURES
You may plot the histograms within their
own figure or within their own subplot of
the same figure
"""
for pred in predictors:
    plt.hist(X[:,pred])
    plt.show()


```







```
[73]: """ TODO
Create a histogram or barplot for the counts
for each target class
"""
plt.hist(y)
plt.show()
```



2.4 BASIC BOXPLOTS OF FEATURES

Boxplots or box-and-whisker plots are used to obtain a perspective of the distribution of the data. The box within the figure displays the 25th percentile (Q1), the median, and the 75th percentile (Q3) of the data. The range between the 75th percentile value and the 25th percentile value is the interquartile range ($IQR = Q3 - Q1$). The end of bottom line is $Q1 - 1.5 * IQR$. The end of top line is $Q3 + 1.5 * IQR$. Anything beyond the lines, the circles, are suggested outliers.


<



>


One can use the `boxplot(data_values, labels=[name])` to generate a boxplot. `data_values` would be the set of observed values for a particular feature and `labels` should be provided as a list, with the name of the feature, in place of `name`.

```
[66]: """ TODO
      BOXPLOTS OF THE CHOSEN PREDICTOR FEATURES
      You may place the boxplots within their
      own figure or within their own subplot of
      the same figure
      """
      print(len(X[0]))
      for pred in predictors:
          plt.boxplot(X[:,pred])
          plt.show()
```

30







2.5 DESCRIPTIVE STATISTICS

```
[13]: # Simply run this cell
      """
      Create a separate variable of the data from the
      predictors
      """
      Xpreds = X[:, predictors]

      """
      Check if any values are NaN (not a number)
      """
      np.any(np.isnan(X[:, predictors]))
```

[13]: False

```
[92]: """ TODO
      Compute the following descriptive statistics of the
      features ignoring NaN values, using numpy:
      mean, median, standard deviation, min, and max

      Make sure to compute the statistics of the columns
      of X (i.e. of each feature). You can specify this
      by setting axis=0 for each of the functions

      Compute and print the results
      """
```

```

for preds in range(4):
    mean = np.mean(Xpreds[:,preds], axis = 0)
    print("The mean of", pred_names[preds], ":", mean)
    median = np.median(Xpreds[:,preds], axis = 0)
    print("The median of", pred_names[preds], ":", median)
    standard_deviation = np.std(Xpreds[:,preds], axis = 0)
    print("The standard deviation of", pred_names[preds], ":",
    →standard_deviation)
    min = np.min(Xpreds[:,preds], axis = 0)
    print("The min of", pred_names[preds], ":", min)
    max = np.max(Xpreds[:,preds], axis = 0)
    print("The max of", pred_names[preds], ":", max)
    print("\n")

```

The mean of mean radius : 14.127291739894552
 The median of mean radius : 13.37
 The standard deviation of mean radius : 3.520950760711062
 The min of mean radius : 6.981
 The max of mean radius : 28.11

The mean of mean area : 654.8891036906855
 The median of mean area : 551.1
 The standard deviation of mean area : 351.60475406323
 The min of mean area : 143.5
 The max of mean area : 2501.0

The mean of mean concavity : 0.0887993158172232
 The median of mean concavity : 0.06154
 The standard deviation of mean concavity : 0.07964972534603185
 The min of mean concavity : 0.0
 The max of mean concavity : 0.4268

The mean of mean symmetry : 0.18116186291739894
 The median of mean symmetry : 0.1792
 The standard deviation of mean symmetry : 0.027390180864268532
 The min of mean symmetry : 0.106
 The max of mean symmetry : 0.304

2.6 FEATURE CORRELATIONS

It's useful to know the correlation between various features, as well as each feature and the predicted label. Feature correlation is useful for feature selection and understanding the

relationship between multiple variables within a dataset. Correlation is either positive, negative, or zero. When two features increase simultaneously, they are positively correlated. When one feature increases while the other decreases, the features are negatively correlated. Zero correlation is when there is no relationship between the features. Correlation is on the range -1 (perfect negative correlation) and 1 (perfect positive correlation).

We can construct scatter plots of one feature versus another to observe linear or nonlinear relationships.

Complete the following set of scatter plots:

<

>

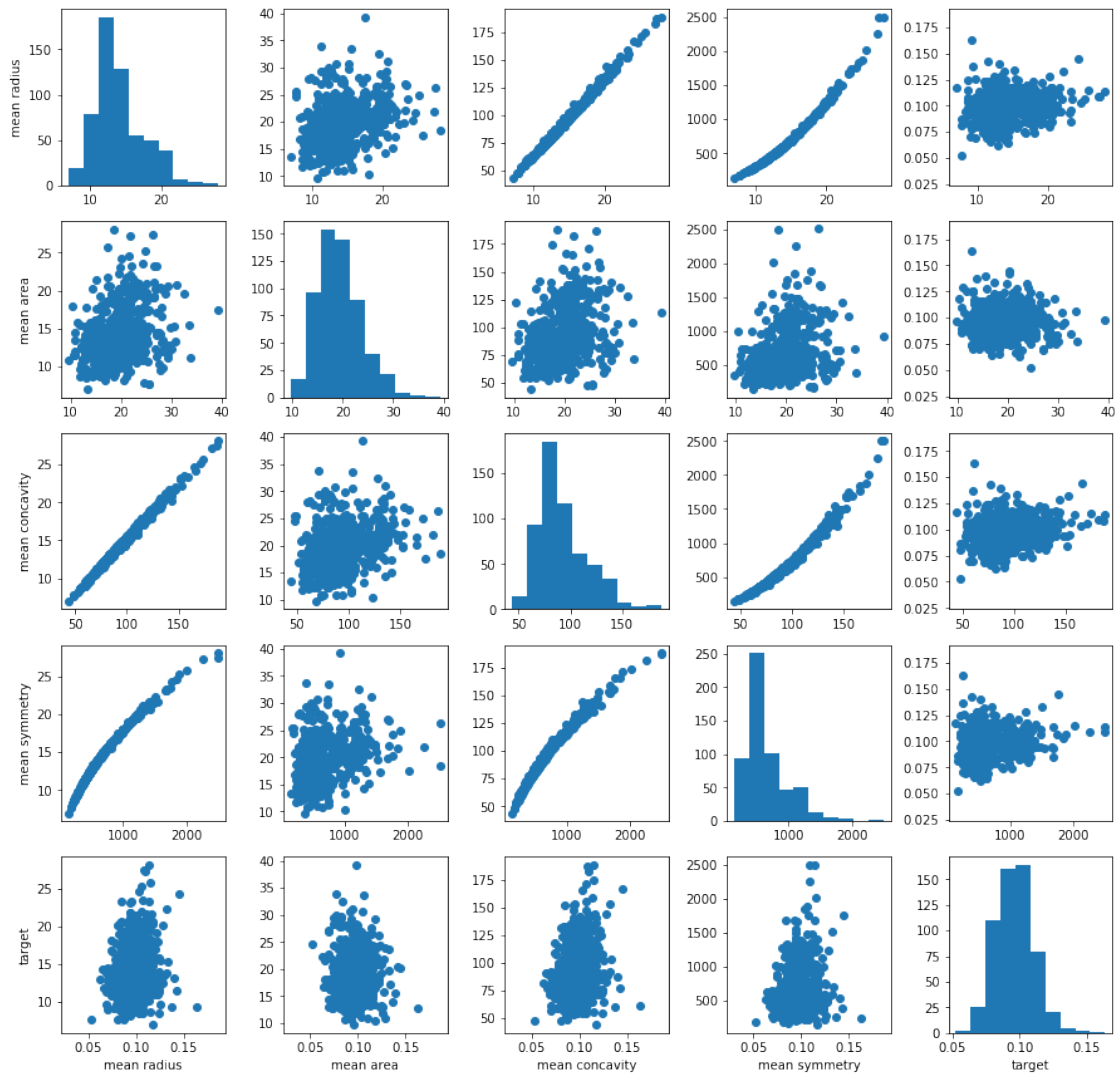
```
[68]: """
Using the scatter plot function, construct plots depicting the
correlation between all pairings of the selected predictor features
and between all predictors and the determined target.
The figure will contain r by r subplots, where r = npredictors + 1.
Where subplot(i,j) is a scatter plot of the feature i versus feature j.
When i == j, plot the histogram of feature i instead of a scatter plot.
We are also interested in the correlation between each of the features
and the target classification, thus we will combine the predictors matrix
and the target vector into one large matrix for convenience.
"""

# Append the y to the end of the matrix of predictors
Xycombo = np.append(Xpreds, y.reshape(-1, 1), axis=1)
# Append the name 'target' to the end of the list of predictor names
Xycolnames = pred_names + ['target']

# Create the scatter plots
fig, axs = plt.subplots(npredictors+1, npredictors+1, figsize=(15, 15))
fig.subplots_adjust(wspace=.35)
for f1 in range(npredictors+1):
    for f2 in range(npredictors+1):

        # v TODO -----
        if f1 == f2:
            axs[f1, f2].hist(X[:, [f1]])
        else:
            axs[f1, f2].scatter(X[:, [f1]], X[:, [f2]])
        # ^ TODO -----

        # include labels only when necessary
        if f1 == npredictors:
            axs[f1, f2].set_xlabel(Xycolnames[f2])
        if f2 == 0:
            axs[f1, f2].set_ylabel(Xycolnames[f1])
```



2.7 IMAGES AND COLORMAPS

Create a colormap plot of the correlations between the all the predictors and the target

```
[107]: """
Generate a figure that plots the a correlation matrix
as a colormap.
PARAMS:
    corrs: matrix of correlations between the features
    varnames: list of the names of each of the features
            (e.g. the column names)
"""
def correlationmap(corrs, varnames):
    nvars = corrs.shape[0]
```

```

# create the figure and plot the correlation matrix
fig, ax = plt.subplots()
im = ax.imshow(corrs, cmap='RdBu', vmin=-1, vmax=1)
cbar = ax.figure.colorbar(im, ax=ax)
cbar.ax.set_ylabel("Pearson Correlation", rotation=-90, va="bottom")

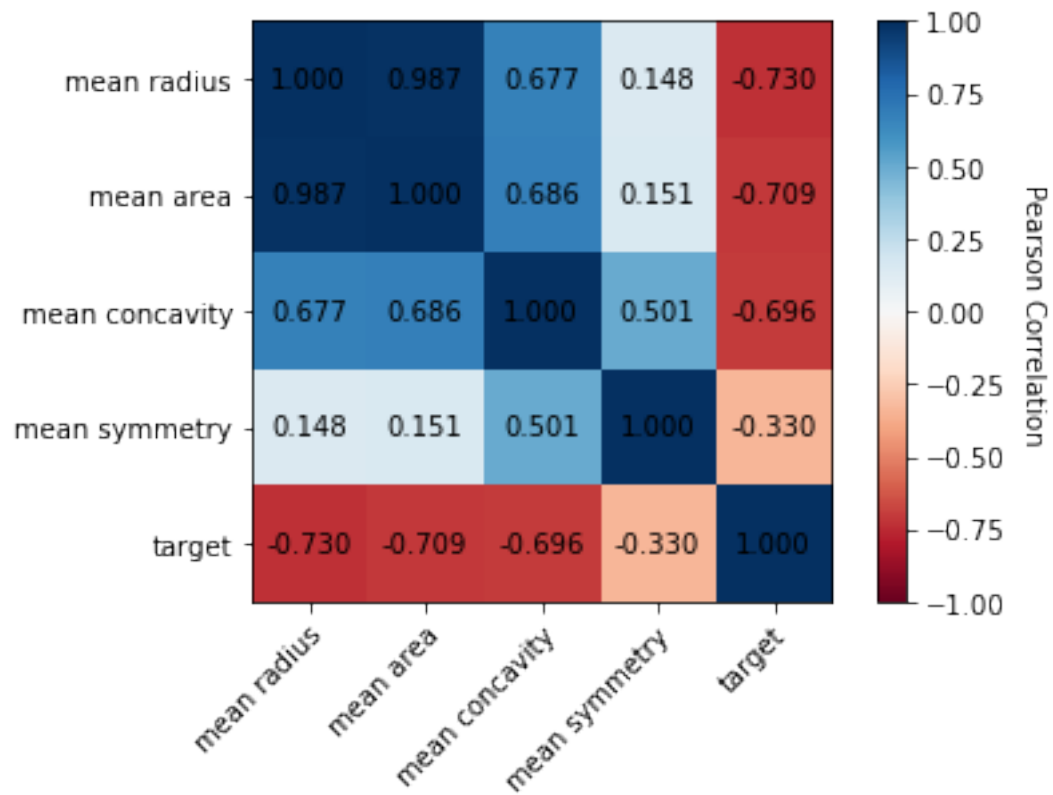
# Specify the row and column ticks and labels for the figure
ax.set_xticks(range(nvars))
ax.set_yticks(range(nvars))
ax.set_xticklabels(varnames)
ax.set_yticklabels(varnames)

# Rotate the tick labels and set their alignment.
plt.setp(ax.get_xticklabels(), rotation=45, ha="right",
rotation_mode="anchor")

# Loop over data dimensions and create text annotations.
for i in range(nvars):
    for j in range(nvars):
        text = ax.text(j, i, "%.3f" % corrs[i, j],
                        ha="center", va="center", color="k")
# END DEF correlationmap -----

"""
Compute the Pearson correlation between the columns of Xycombo using
the numpy function corrcoef(). The corrcoef() function performs the
the pairwise correlation on the rows of a matrix, thus you will need to
transpose the input.
"""
Xytranspose = np.matrix.transpose(Xycombo)
Xycorrs = np.corrcoef(Xytranspose)
""" TODO
Call the function defined above, correlationmap(), to generate a
colormap plot of the correlations between columns of the Xycombo matrix.
Mke sure to read any provided code
"""
correlationmap(Xycorrs, Xycolnames)

```

[]: